

Architecture Overview

This deployment strategy implements a **modern, serverless, cloud-native architecture** on AWS that uses AWS managed services to minimize infrastructure overhead while maximizing scalability, reliability, and security.

1. API Deployment:

- **Container-based:** Flask API packaged in Docker and stored in Amazon ECR
- **Serverless compute:** Runs on Amazon ECS with AWS Fargate (no server management)
- **Auto-scaling:** Application Load Balancer distributes traffic and scales containers based on demand

2. Database Management:

- **Fully managed:** Amazon RDS PostgreSQL with automated backups, patching, and maintenance
- **High availability:** Multi-AZ deployment for instant failover protection
- **Security-first:** Isolated in VPC with credentials managed by AWS Secrets Manager
- **Performance:** Built-in monitoring and optimization capabilities

3. Data Ingestion:

- **Scheduled processing:** Amazon EventBridge triggers data ingestion jobs
- **Serverless execution:** AWS Lambda functions process weather data from S3
- **Error resilience:** Dead Letter Queues ensure no data loss during failures
- **Cost-effective:** Pay-per-execution model with automatic scaling

Supporting Infrastructure

- **CI/CD Pipeline:** Automated deployment using AWS CodePipeline and CodeBuild
- **Monitoring:** CloudWatch and X-Ray provide comprehensive observability
- **Security:** AWS WAF protection with VPC network isolation
- **Cost Optimization:** Serverless and managed services reduce operational costs

Key Benefits

- **Scalability:** Automatically handles traffic spikes and data processing loads
- **Reliability:** Multi-AZ redundancy and automated failover mechanisms
- **Security:** Defense-in-depth approach with encryption and network isolation
- **Cost-Efficiency:** Pay-only-for-what-you-use serverless model
- **Operational Excellence:** Minimal infrastructure management with maximum automation